

Tech Design Document – GitHub Actions Pipeline Dashboard

This document explains the system in an engineer-friendly **and non-technical** way. It covers how the dashboard works, what talks to what, which APIs exist, what we store in the database, and how the user interface is organized.

1) High-level architecture

Purpose in one line

A small web app that **collects GitHub Actions run data**, stores it in a database, **shows live metrics** (success rate, average build time, last build status), and **alerts on failures** (Slack/email).

Main building blocks

1. **Frontend (React + Vite, served by Nginx in Docker)**

The web page you open in the browser. It shows cards, charts, and a table of recent runs, and calls the backend's REST APIs.

2. **Backend (Node.js + Express, Docker)**

The application server. It:

- Polls GitHub's REST API for workflow runs at a fixed interval.
- Calculates metrics (success/failure rate, average build time, latest run status).
- Stores runs in **PostgreSQL**.
- Exposes REST APIs that the frontend calls.
- Sends **Slack** alerts when a run fails.

3. **Database (PostgreSQL, Docker)**

Stores the raw runs and basic computed fields like duration (in seconds). This enables historical views and reliable metrics.

4. **Alerting (Slack webhook)**

When the backend sees a failed run, it posts a concise message to a Slack channel. Optional email integration can be added similarly.

5. **Containerization (Docker Compose)**

Brings up Postgres, backend, and frontend in a single command. Each component is isolated and can be restarted independently.

How data flows end-to-end

1. A GitHub Action workflow finishes (success/failure).
2. On a schedule (e.g., every 60 seconds), the backend **polls GitHub** for the latest runs.
3. The backend **upserts** those runs into Postgres (insert if new, update if already seen).
4. If any run has **conclusion = failure**, the backend posts an **alert to Slack**.
5. The frontend calls the backend **metrics and runs APIs** and renders the dashboard.

Why this architecture

- **Simple & reliable:** polling avoids webhook setup hurdles; Slack alerts provide immediate visibility.
 - **Portable:** all components are Dockerized and can run on a laptop, VM, or server.
 - **Extensible:** can add Jenkins, CircleCI, or multiple GitHub repos later with minimal changes.
-

2) API structure (routes & sample responses)

Base URL: /api

Notes: • All responses are JSON. • Errors follow { "error": "message" } with appropriate HTTP status codes.

2.1 Health

GET /api/health

Checks that the API is up.

```
{ "ok": true }
```

2.2 Metrics

GET /api/metrics?windowHours=24

Returns summary metrics for the selected time window.

```
{
  "windowHours": 24,
  "success": 42,
  "failure": 8,
  "completed": 50,
  "successRate": 0.84,
  "averageBuildTimeSec": 315,
  "latestRun": {
    "id": 1234567890,

```

```

    "status": "completed",
    "conclusion": "success",
    "name": "CI",
    "branch": "main",
    "event": "push",
    "actor": "octocat",
    "created_at": "2025-08-21T10:15:04.000Z",
    "run_started_at": "2025-08-21T10:15:05.000Z",
    "updated_at": "2025-08-21T10:20:20.000Z",
    "html_url": "https://github.com/org/repo/actions/runs/1234567890",
    "duration_seconds": 315
  }
}

```

2.3 Runs (paginated)

GET /api/runs?page=1&pageSize=20

Lists the most recent runs.

```

{
  "page": 1,
  "pageSize": 20,
  "items": [
    {
      "id": 1234567890,
      "status": "completed",
      "conclusion": "failure",
      "name": "CI",
      "branch": "feature-x",
      "event": "push",
      "actor": "jane",
      "created_at": "2025-08-21T09:03:10.000Z",
      "run_started_at": "2025-08-21T09:03:12.000Z",
      "updated_at": "2025-08-21T09:05:56.000Z",
      "html_url": "https://github.com/org/repo/actions/runs/1234567890",
      "duration_seconds": 164
    }
  ]
}

```

2.4 Latest run (shortcut)

GET /api/latest-run

Returns a single most recent run.

```

{
  "id": 1234567890,
  "status": "completed",
  "conclusion": "success",
  "name": "CI",

```

```
"branch": "main",
"event": "push",
"actor": "octocat",
"created_at": "2025-08-21T10:15:04.000Z",
"run_started_at": "2025-08-21T10:15:05.000Z",
"updated_at": "2025-08-21T10:20:20.000Z",
"html_url": "https://github.com/org/repo/actions/runs/1234567890",
"duration_seconds": 315
}
```

2.5 Manual poll

POST /api/poll

Forces an immediate GitHub poll (useful during demos).

```
{ "ingested": 37 }
```

2.6 Errors (examples)

- 401 Unauthorized from GitHub → the backend will log a clear message like “*GitHub token missing or invalid*” and respond with an error if the endpoint proxies anything.
- 500 Internal Server Error → generic unexpected error (e.g., DB not reachable).

3) Database schema

3.1 Tables

Table: runs (one row per GitHub Actions workflow run)

Column	Type	Notes
id	BIGINT (PK)	GitHub Actions run ID (globally unique per run)
status	TEXT	queued, in_progress, completed
conclusion	TEXT	success, failure, cancelled, skipped, null (if not done)
name	TEXT	Workflow name
event	TEXT	e.g., push, pull_request
branch	TEXT	Derived from head_branch
actor	TEXT	GitHub username who triggered it
created_at	TIMESTAMPZ	When the run record was created by GitHub
run_started_at	TIMESTAMPZ	When the run actually began

Column	Type	Notes
updated_at	TIMESTAMPTZ	Last update time (finish time for completed runs)
html_url	TEXT	Link to the run in the GitHub UI
run_attempt	INT	1 for first attempt, increments on retries
duration_seconds	INT	Calculated: updated_at - run_started_at in seconds

Indexes: - idx_runs_created_at on (created_at DESC) → fast latest queries. - idx_runs_branch on (branch) → filter by branch quickly. - idx_runs_conclusion on (conclusion) → reporting by outcome.

Initialization DDL (excerpt):

```
CREATE TABLE IF NOT EXISTS runs (
  id BIGINT PRIMARY KEY,
  status TEXT NOT NULL,
  conclusion TEXT,
  name TEXT,
  event TEXT,
  branch TEXT,
  actor TEXT,
  created_at TIMESTAMPTZ,
  run_started_at TIMESTAMPTZ,
  updated_at TIMESTAMPTZ,
  html_url TEXT,
  run_attempt INT,
  duration_seconds INT
);
```

```
CREATE INDEX IF NOT EXISTS idx_runs_created_at ON runs(created_at DESC);
CREATE INDEX IF NOT EXISTS idx_runs_branch ON runs(branch);
CREATE INDEX IF NOT EXISTS idx_runs_conclusion ON runs(conclusion);
```

3.2 Data lifecycle

- **Insert/Update:** on each poll, new runs are inserted; existing runs are updated (retry count, conclusion, timestamps, duration).
 - **Retention (optional):** you may periodically archive old rows to S3 or prune after N days.
 - **Aggregation:** metrics are computed at query time (no separate summary tables needed for this scope).
-


4) UI layout (explanation)


Overview


The UI focuses on **clarity and at-a-glance health**.

Header

- Title: *GitHub Actions Dashboard*
- Subtext: repository/owner and selected time window

Metric cards (top row) -  **Success Rate** — percentage of completed runs that ended in success within the window.

-  **Average Build Time (s)** — average of duration_seconds across completed runs in the window.

-  **Last Build Status** — status and branch of the most recent run; includes a link to GitHub. - **Window selector** — dropdown to choose 1h/6h/12h/24h/48h/72h/7d. - **Actions** — a **Poll Now** button to trigger immediate refresh via /api/poll.

Charts & tables (main area) - **Bar chart: Success vs Failure** — helps spot trends (e.g., sudden spike in failures).

- **Recent Runs table** — sortable columns: Run ID, Name, Branch, Status, Conclusion, Duration (s), Link.

- Clicking the **Logs** link opens the run in GitHub for detailed logs.

Alerting feedback (optional panel) - Shows latest Slack alert timestamp (“Last failure alert sent ...”).

Empty/Loading states - Clear copy like “No runs in the selected window yet” and skeleton loaders during fetch.

Responsiveness & accessibility - Cards stack on small screens; table scrolls horizontally if needed.

- Color + text indicators (not color-only) for status to support color-blind users.

Interaction model

- The UI polls **only the backend**; the backend in turn polls GitHub.
- Time window changes trigger a fresh /api/metrics fetch.
- “Poll Now” calls /api/poll and then reloads metrics and runs.

Non-functional notes (quick)

- **Security:** GitHub PAT stored as environment variable; never committed. Backend rejects startup if token missing. CORS restricted to the frontend origin

(configurable).

- **Performance:** pagination for runs; indexes on time/branch/outcome; polling interval configurable (e.g., 60s).
 - **Reliability:** if GitHub is temporarily unavailable, backend logs the error and retries on next interval; Postgres connection pooled.
 - **Extensibility:** add org-wide or multi-repo monitoring by storing owner, repo columns and looping over a configured list.
-